

NeuroGrad Repository Analysis

Architecture and Design

- **Repository Structure:** The project is organized into a Python package with clear submodules ¹ ² :
 - `neurograd/tensor.py` – defines the core **Tensor** class (primary data structure with autograd support).
 - `neurograd/functions/` – differentiable math operations (e.g. arithmetic ops, activations, conv, reductions), each implemented as a subclass of a base **Function** ¹ .
 - `neurograd/amp/` – mixed-precision support (Autocast context manager and gradient scaler for FP16 training) ³ .
 - `neurograd/nn/` – neural network components: layers (Linear, Conv2D, Pooling, Dropout, etc.), loss functions, evaluation metrics, and the **Module** base class for building models ⁴ .
 - `neurograd/optim/` – optimization algorithms (SGD, Adam, RMSprop implementations) ⁵ .
 - `neurograd/utils/` – utilities for data loading, gradient checking, graph visualization, device management, etc. ⁶ .
- **Design Principles:** NeuroGrad follows a design similar to PyTorch's **dynamic computation graph** and module system. Each operation is tracked in a runtime graph for automatic backpropagation ⁷ . The Tensor/Function design pattern means every operation creates a Function node that knows how to `forward` compute results and `backward` propagate gradients ⁸ ⁹ . Neural network layers are built as subclasses of `Module` with an overridable `forward()` method and automatic parameter registration ¹⁰ ¹¹ . The framework uses a **backend abstraction**: a unified `xp` alias points to NumPy (CPU) or CuPy (GPU), so computations can run on either device seamlessly ¹² . This modular, object-oriented design (core Tensor + Function for autodiff, Module for layers) makes the code organized and extensible.

Model and Training Capabilities

- **Supported Model Types:** NeuroGrad supports **feed-forward neural networks** including fully-connected networks (MLPs) and convolutional neural networks (CNNs). It provides building blocks like `Linear` layers for dense networks and `Conv2D` / `Pool2D` layers for convolutional models ¹³ . Complex models can be constructed using the `Module` API or the provided `Sequential` container (for stacking layers in order) ¹⁴ ¹⁵ . For example, an `MLP` class is included that quickly assembles a multi-layer perceptron from a list of layer sizes ¹⁶ , and a CNN can be built by combining `Conv2D`, pooling, and `Linear` layers (as shown in the examples).
- **Layer Features:** The framework's layers come with common enhancements. The `Linear` layer supports optional **activation functions**, **dropout**, and **batch normalization** built-in ¹⁷ ¹⁸ . Similarly, `Conv2D` layers allow specifying activation, dropout, and batchnorm on the outputs ¹⁹ .

²⁰ . These layers use proper weight initialization schemes (He or Xavier initialization for weights, zeros or others for biases) to promote stable training ²¹ ²² . Activation functions provided include ReLU, Sigmoid, Tanh, LeakyReLU, Softmax, etc., which can be used standalone or by name in layer definitions ²³ . Common **loss functions** are implemented (MSE, RMSE, MAE for regression; Binary Cross-Entropy and Categorical Cross-Entropy for classification) ²⁴ , and **metrics** like accuracy, F1-score, etc., help evaluate model performance ²⁵ .

- **Training Process:** Model training in NeuroGrad follows the standard **forward-backward optimization loop** controlled by the user (similar to PyTorch's training pattern). The library provides a `Dataset` and `DataLoader` utility to iterate over mini-batches of data ²⁶ ²⁷ . Users compose a training loop where for each batch, the model (a `Module`) is called to get predictions, a loss is computed, and `loss.backward()` is invoked to compute gradients. Then an optimizer (SGD/Adam/RMSprop) is used to update parameters ²⁸ ²⁹ . NeuroGrad's optimizers have an API analogous to PyTorch's: e.g. `optimizer.zero_grad()` to reset gradients and `optimizer.step()` to apply updates ³⁰ . An example from the documentation trains an MLP on a dataset for multiple epochs, printing loss and accuracy each epoch ²⁸ ²⁹ . This explicit loop design makes it clear how training proceeds and allows easy customization (learning rate schedules, custom logging, etc.).
- **Optimizations for Training:** Despite being written in pure Python, NeuroGrad includes some advanced training optimizations. It supports **automatic mixed-precision (AMP) training**: by using `ng.amp.autocast` and `GradScaler`, users can train with float16 where safe, gaining speed and memory benefits ³¹ ³² . The example CNN training code shows wrapping the forward pass in an autocast context and scaling the loss before backward, then un-scaling during optimizer step ³² ³³ . This yields a reported **1.5-2× speedup and ~40-50% memory reduction** without sacrificing accuracy ³⁴ . Furthermore, **GPU acceleration** is available: if CuPy is installed and a CUDA device is detected, the `xp` backend uses CuPy so that all tensor operations and linear algebra run on the GPU ³⁵ . The optimizers are also optimized for GPU usage – they use fused elementwise operations via CuPy's JIT when possible to update weights with minimal Python overhead ³⁶ . Overall, models can be trained on CPU or GPU, in full or mixed precision, using standard techniques like momentum (in SGD) ³⁷ and adaptive moments (Adam) ³⁸ , which are fully supported by NeuroGrad.

Datasets and Input Formats

- **Dataset Classes:** NeuroGrad does not hard-code specific datasets but provides flexible classes to handle input data. The primary interface is the `Dataset` class, which can wrap any pair of arrays (features `X` and labels `y`) into an iterable dataset object ³⁹ . It stores `X` and `y` internally as NeuroGrad Tensors and supports indexing and length queries ⁴⁰ . A corresponding `DataLoader` can then be used to batch and shuffle data – it takes a `Dataset` and yields mini-batches of a given size, returning batched tensors for inputs and targets ²⁶ ⁴¹ . This design is similar to PyTorch's `DataLoader`, making it easy to plug in numpy data or custom data sources.
- **Supported Input Formats:** The framework is quite general in terms of input data. Since the `Tensor` constructor will convert inputs using `xp.array`, any data type compatible with NumPy/CuPy (Python lists, NumPy ndarrays, etc.) can be used as input. For image data, NeuroGrad includes a convenient `ImageFolder` dataset class that mimics PyTorch's `ImageFolder`: you point it to a

directory, and it loads all images from subfolders as class-labeled samples ⁴² ⁴³. It supports common image file formats (`.png`, `.jpg`, `.bmp`, etc.) ⁴⁴ and can apply transformations like resizing to a specified shape, converting to grayscale or RGB, and normalizing pixel values (by default it scales pixels to `[0, 1]` floats by dividing by 255) ⁴⁵. The loaded images are automatically wrapped in Tensors of type `float32` (or another dtype if specified) and labels are converted to integer class indices ⁴⁶ ⁴³. This means out-of-the-box one can train on image datasets organized in folders, with optional preprocessing.

- **Data Preprocessing:** NeuroGrad's dataset utilities handle basic preprocessing steps like shuffling data and normalizing images, but more complex preprocessing (feature scaling, augmentation, etc.) is left to the user or provided as hooks. For instance, `ImageFolder` accepts an `img_transform` callback to apply custom numpy transformations on each image, and a `target_transform` for labels ⁴⁷ ⁴⁸. In practice, users prepare data as needed (e.g., flattening images or one-hot encoding labels) before creating a `Dataset`. Notably, classification labels are often expected as one-hot vectors when using the built-in cross-entropy loss – the **CategoricalCrossEntropy** loss multiplies the target vector with `log(predictions)` internally ⁴⁹ ⁵⁰. (If `from_logits=True` is used, the loss will apply a softmax and still assume `y_true` is one-hot ⁵¹.) This implies that for multi-class classification tasks, the user should convert label indices to one-hot vectors prior to training (as demonstrated in the example notebooks).
- **Example Usage:** The repository's examples illustrate the data handling. In a provided Jupyter notebook, a LeNet-5 CNN is trained on the **MNIST** dataset (handwritten digits) – the code downloads the MNIST images and labels, then uses NeuroGrad to construct a `Dataset` and `DataLoader` for training ⁵². The images (28×28 grayscale) are padded to 32×32 and normalized, and the labels are converted to one-hot encoding in that example. This confirms that **MNIST** (a standard benchmark) is used as a proving ground for the framework. While MNIST is not hardcoded into NeuroGrad itself, the framework readily supports it and similarly structured datasets. In general, any numeric data in NumPy/CuPy format can serve as input, making NeuroGrad applicable to a wide range of tasks (vision, basic NLP with sequences as tensors, etc., though sequence models would need additional support as noted below).

Key Algorithms and Innovations

- **Automatic Differentiation Engine:** A core innovation of NeuroGrad is its from-scratch implementation of reverse-mode autodiff. The **Tensor** class lies at the heart of this – it holds a value (data array), a `grad` attribute for the gradient, and a reference to a `grad_fn` (the Function that produced it) ⁵³ ⁵⁴. As operations are applied to Tensors, NeuroGrad creates a computational graph of Function objects. Each **Function** subclass implements `forward` (to compute the output) and `backward` (to compute gradients of its inputs given the gradient of its output) ⁸ ⁹. During backpropagation (`Tensor.backward()`), the engine performs a **topological sort** of the graph and accumulates gradients in reverse order ⁵⁵ ⁵⁶. This ensures each operation's backward is called only after all of its dependents' gradients are known. The framework also includes safeguards: for example, it checks that `.backward()` is only called on scalar outputs or with a gradient passed in ⁵⁷, and it avoids repeated gradient computation on shared subgraphs by marking visited nodes ⁵⁸. After computing grads, by default it **frees the graph references** (setting `grad_fn=None` on intermediates) to avoid memory leaks unless `retain_graph=True` is specified ⁵⁹. Overall, the

autodiff system is quite sophisticated for a lightweight library – it handles arbitrary graph structures (including branchings and broadcast operations) and supports second-order gradients if needed (since the `Function` classes themselves are Modules and can be differentiated through, as long as you keep the graph).

- **Custom Gradient Implementations:** NeuroGrad’s source code provides manual gradient formulas for each operation, which is a key educational aspect. Many basic ops (add, subtract, multiply, etc.) have straightforward backward passes (e.g., propagating the gradient or applying broadcasting rules). More complex ops illustrate clever techniques: for example, the **Pad** operation’s backward simply slices out the padded regions from the gradient ⁶⁰ ⁶¹. The implementation of **convolution** is particularly notable. Convolution forward is done by an **im2col + matrix multiplication** approach – using NumPy’s `stride_tricks` to create a sliding window view of the input and then reshaping and multiplying to apply filters ⁶². Instead of deriving a monolithic backward for Conv2D, NeuroGrad composes it from simpler operations (transpose, reshape, matmul, etc.) whose gradients are known. Critical to this is the `SlidingWindowView` Function, which underpins both convolution and pooling. Its **backward** is implemented to efficiently scatter the gradients back to the input: it maintains a zero-initialized buffer of the input’s shape and adds the gradient of each window into this buffer using an overlapping view ⁶³ ⁶⁴. This approach avoids allocating giant intermediate tensors – the sliding window’s backward reuses a preallocated buffer and accumulates in place, which is memory-efficient. Such implementations (using NumPy/CuPy cleverly) demonstrate the framework’s innovative balance between clarity and performance.

- **Mixed Precision and AMP:** NeuroGrad stands out among pure-Python frameworks by providing **automatic mixed-precision training** support. It introduces an `autocast` context manager and a `GradScaler` similar to PyTorch’s, which required solving how to integrate with the autograd system. The solution was to use a utility that conditionally casts tensors to float16 within `Function.__call__` before the forward computation, except for operations known to be unsafe in lower precision ⁶⁵ ⁶⁶. This means whenever `autocast(enabled=True)` is in effect, each operation checks if it should run in FP16 or FP32. For example, matrix multiplications and convolution might run in FP16 for speed, while losses or certain reductions remain in FP32 for numeric stability. After the forward pass, `GradScaler.scale(loss)` multiplies the loss by a scale factor and on `.step()` it unscales gradients and adjusts the factor ⁶⁷. All of this is done with a **PyTorch-compatible API** (the context and scaler have the same interface) ³² ³³, making it easy for users familiar with AMP to adopt. The effective result is significantly faster training on GPUs (as noted, up to ~2x faster) with much lower memory usage, achieved without requiring the user to manually manage any casting. This is a notable feature since many educational frameworks stick to full precision; NeuroGrad shows an innovative engineering effort to bridge that gap.

- **Optimizer Fusions:** The optimization algorithms themselves incorporate some low-level optimizations. When running on GPU (CuPy backend), NeuroGrad uses **fused CUDA kernels** for updating parameters. In the SGD implementation, for instance, the momentum update and weight update are combined into single fused element-wise operations using CuPy’s `@cupy.fuse` decorator ³⁶. This JIT-compiles the update step (`momentum = beta*momentum + (1-beta)*grad` and `param -= lr * momentum`) into one GPU kernel, reducing Python overhead in the inner loop. Similarly, Adam’s update of first and second moment estimates and the bias-corrected weight update are fused into a few kernels ⁶⁸. These kinds of optimizations are quite advanced for a pure Python framework – they indicate that while NeuroGrad is educational, it also

considers performance best practices (especially for GPU training). Additionally, the optimizers support **weight decay** (L2 regularization) simply by in-place adding the scaled weights to the grad before the update ³⁶, which is convenient for users to enable.

- **Visualization and Debugging Tools:** NeuroGrad includes features to help understand and debug models. There is a **graph visualization** utility that can create a visual graph of the computation. For any output tensor, calling `tensor.visualize_graph()` returns a Matplotlib figure illustrating nodes (tensors and functions) and their connections ⁶⁹ ⁷⁰. This is helpful for seeing how a complex model's graph is structured (for example, it would show a chain of Linear and activation operations in an MLP graph). The `graph.py` utility also allows printing a textual graph and computing stats like number of nodes and graph depth ⁷¹. Another helpful tool is the **gradient checker**: the `utils/grad_check.py` module provides a `gradient_check()` function that compares the analytic gradients from NeuroGrad's autograd to numerical gradients (using small finite differences) for any model and loss ⁷². This is extremely useful when implementing new Function classes – you can quickly validate if your backward is correct. These innovations underscore the educational focus: a user can both visualize what the model is doing internally and verify that gradients are computed properly, which aids learning and development of new features.

Usability and Documentation

- **Documentation Quality:** The NeuroGrad repository is accompanied by thorough documentation that makes it approachable. The **README.md** on GitHub provides a clear overview of the framework's purpose and capabilities ⁷³. It highlights key features and lists all major components (autodiff, layers, optimizers, etc.) in bullet form ⁷⁴ ⁷⁵ for quick scanning. The README includes step-by-step **installation instructions** (both via PyPI and from source) and a "Quick Start" code example to get new users started quickly ⁷⁶ ⁷⁷. There are also **usage examples** demonstrating core functionality: creating tensors and performing operations ⁷⁸ ⁷⁹, as well as a complete training loop example for a simple MLP model training on a dataset ²⁸ ²⁹. Throughout the README, code snippets are provided alongside explanations, which greatly lowers the barrier to entry for new users.
- **API Familiarity:** The framework's API is intentionally designed to mirror the style of popular libraries like PyTorch, which improves usability for those with prior deep learning experience. For instance, creating a tensor via `ng.Tensor(data, requires_grad=True)` and calling `tensor.backward()` is analogous to PyTorch's `torch.tensor` and `.backward()` usage ⁷⁸ ⁷⁹. The `Module` class and layer interface (with methods like `.parameters()` and `.zero_grad()`) behave similarly to `torch.nn.Module` ⁸⁰ ⁸¹. Even the mixed-precision API (`autocast`, `GradScaler`) is made compatible with PyTorch's conventions ³². This consistency means users can transfer knowledge easily and the learning curve is gentle. In addition, the naming of classes and functions is very transparent (e.g., `Linear`, `Conv2D`, `SGD`, `accuracy_score`), and the repository contains docstrings and comments for further clarification. Key design concepts are explained in the documentation; for example, the CLAUDE.md file in the repo (a developer guide) explicitly describes the architecture and design patterns, which helps advanced users understand the library's internals ⁸² ⁸³.

- **Ease of Getting Started:** To try NeuroGrad, one can install it via `pip install neurograd` (it even provides extras like `[gpu]` for GPU support and `[all]` for full features) ⁸⁴. After installation, the user can follow the code examples to create models. The README's "**Quick Start**" shows basic tensor math and autodiff, which reinforces how to use the Tensor API. For example, the user can replicate a simple computation: create tensors, do `z = x @ y + x.sin()`, then call `z.sum().backward()` and inspect `x.grad` ⁸⁵ ⁸⁶. This guided introduction helps verify that everything is working as expected. The README then provides a complete **training script example**, which newcomers can use as a template for training their own models ²⁸ ²⁹. There's also an example of using the new mixed-precision features in a training loop, which is helpful for users interested in performance tuning ³². Overall, the combination of familiar interfaces, example-driven documentation, and the educational tone (the project explicitly targets learning and prototyping ⁸⁷) makes the framework very user-friendly.
- **Support and Extensibility:** NeuroGrad's documentation encourages contributions and provides guidance for extending the framework. The README outlines a roadmap of upcoming features and invites bug reports or feature discussions via GitHub issues ⁸⁸ ⁸⁹. For those looking to extend NeuroGrad, there are instructions (in CLAUDE.md and comments) on how to add a new Function or layer – e.g., it suggests steps like inheriting from `Function`, implementing `forward`/`backward`, and using the gradient checker to test it ⁹⁰ ⁹¹. The presence of a test suite (with `pytest` instructions given) and continuous integration of concepts like gradient checking means developers can gain confidence in any new code they write. From a usability perspective, this transparency and welcoming of extension mean that advanced users can treat NeuroGrad not just as a black-box tool, but as a learning platform to experiment with new ideas (such as trying out a novel activation or optimizer). This dual nature – easy for beginners to pick up, yet open for experts to tinker with – is well supported by the documentation and code quality.

Limitations and TODOs

- **Current Limitations:** While NeuroGrad is a powerful educational tool, it has some **limitations** in its current state. Firstly, it does not natively support **recurrent neural networks** or sequence models – all provided layers are feed-forward, so architectures like RNNs, LSTMs, or GRUs are not available yet. Similarly, **computer vision** tasks are supported through Conv2D, but more complex layers (e.g. transposed convolutions for generative models) are not present. Another limitation is in model serialization: one can save and load parameter state dictionaries manually ⁹² ⁹³, but there isn't a high-level model save/load utility or ONNX export, etc., which means checkpointing models is a bit manual. There is also no support for **distributed training** – you can use a single GPU or CPU, but multi-GPU training or multi-node training features (as found in more industrial frameworks) are not implemented. In terms of data handling, while basic dataset classes exist, there's no built-in data augmentation pipeline or dataset downloads (users must handle those outside the library). Performance-wise, being pure Python and numpy-based, NeuroGrad may be slower on very large-scale tasks compared to frameworks written in C/CUDA (though CuPy mitigates this on GPU). These constraints are mostly due to the project's young, educational nature and are acknowledged by the developers.
- **Planned Features (Roadmap):** The repository explicitly lists a **roadmap of upcoming features** ⁹⁴, indicating that many of the above limitations are intended to be addressed. Key items in development include:

- **Recurrent Layers** – adding support for RNN, LSTM, GRU layers to handle sequential data ⁹⁵. This will expand the framework's coverage to time-series and language tasks.
- **Additional Optimizers** – implementing more advanced or specialized optimization algorithms like AdaGrad, Nadam, etc. ⁹⁵. These would complement SGD/Adam and give users more choices for training strategies.
- **Model Saving/Loading** – introducing an easier way to save models (e.g. serialization of `Module` objects) and load them back ⁹⁶. Currently, users rely on `state_dict()` and external pickle utilities, so an official method would improve usability for deployment or checkpointing.
- **Distributed Training** – exploring support for distributed computing to train models on multiple GPUs or machines in parallel ⁹⁷. This is a complex feature but is on the radar for future enhancement, which would move NeuroGrad closer to production-grade capabilities.
- **Quantization and Pruning** – tools for model compression, like dynamic quantization (reducing model precision for inference efficiency) and pruning of weights ⁹⁸. These are advanced features usually found in mature frameworks, and their inclusion in the roadmap shows an ambition to grow NeuroGrad's feature set for research experiments.
- **Ongoing Development:** In addition to the roadmap items, the developers likely will continue to refine the existing code. The presence of a test suite and continuous improvements (e.g., the recent addition of mixed-precision support) suggest that NeuroGrad is actively maintained. Some **known areas of improvement** could include optimizing performance (perhaps using numba or more CuPy fusion in the future), expanding the library of activation functions (the framework already demonstrates how to add a custom `Swish` function in the README ⁹⁹ ¹⁰⁰), and improving numerical stability for very deep networks. The project explicitly notes it was built **“from scratch with no AI assistance”** and is meant for learning and prototyping ¹⁰¹. As such, some trade-offs have been made in favor of clarity over efficiency. However, the planned features indicate that many gaps will be filled. Users of NeuroGrad can expect its capabilities to grow, and they are encouraged to contribute (as outlined in the contributing section of the docs). In summary, **current limitations** are primarily missing advanced features rather than flaws in existing components, and the **TODOS** on the roadmap show a clear path toward a more complete deep learning framework in future releases.

1 2 3 4 5 6 13 23 24 28 29 30 31 32 33 34 35 67 69 70 71 72 73 74 75 76 77 78 79 84
85 86 87 88 89 94 95 96 97 98 99 100 101 **README.md**

<https://github.com/b-ionut-r/neurograd/blob/b1de3737c5c20332260885f2ad46ef6b9c05d38a/README.md>

7 10 12 25 82 83 90 91 **CLAUDE.md**

<https://github.com/b-ionut-r/neurograd/blob/b1de3737c5c20332260885f2ad46ef6b9c05d38a/CLAUDE.md>

8 9 65 66 **base.py**

<https://github.com/b-ionut-r/neurograd/blob/b1de3737c5c20332260885f2ad46ef6b9c05d38a/neurograd/functions/base.py>

11 14 15 80 81 92 93 **module.py**

<https://github.com/b-ionut-r/neurograd/blob/b1de3737c5c20332260885f2ad46ef6b9c05d38a/neurograd/nn/module.py>

16 17 18 21 22 **linear.py**

<https://github.com/b-ionut-r/neurograd/blob/b1de3737c5c20332260885f2ad46ef6b9c05d38a/neurograd/nn/layers/linear.py>

19 20 **conv.py**

<https://github.com/b-ionut-r/neurograd/blob/b1de3737c5c20332260885f2ad46ef6b9c05d38a/neurograd/nn/layers/conv.py>

26 27 39 40 41 42 43 44 45 46 47 48 **data.py**

<https://github.com/b-ionut-r/neurograd/blob/b1de3737c5c20332260885f2ad46ef6b9c05d38a/neurograd/utils/data.py>

36 37 **sgd.py**

<https://github.com/b-ionut-r/neurograd/blob/b1de3737c5c20332260885f2ad46ef6b9c05d38a/neurograd/optim/sgd.py>

38 68 **adam.py**

<https://github.com/b-ionut-r/neurograd/blob/b1de3737c5c20332260885f2ad46ef6b9c05d38a/neurograd/optim/adam.py>

49 50 51 **losses.py**

<https://github.com/b-ionut-r/neurograd/blob/b1de3737c5c20332260885f2ad46ef6b9c05d38a/neurograd/nn/losses.py>

52 **lenet5.ipynb**

<https://github.com/b-ionut-r/neurograd/blob/b1de3737c5c20332260885f2ad46ef6b9c05d38a/lenet5.ipynb>

53 54 55 56 57 58 59 **tensor.py**

<https://github.com/b-ionut-r/neurograd/blob/b1de3737c5c20332260885f2ad46ef6b9c05d38a/neurograd/tensor.py>

60 61 63 64 **tensor_ops.py**

https://github.com/b-ionut-r/neurograd/blob/b1de3737c5c20332260885f2ad46ef6b9c05d38a/neurograd/functions/tensor_ops.py

62 **conv.py**

<https://github.com/b-ionut-r/neurograd/blob/b1de3737c5c20332260885f2ad46ef6b9c05d38a/neurograd/functions/conv.py>